

In [1]:

```
# This Python 3 environment comes with many helpful analytics libraries installed
# It is defined by the kaggle/python Docker image: https://github.com/kaggle/docker-python
# For example, here's several helpful packages to load

import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
sns.set_style('darkgrid')

# Input data files are available in the read-only "../input/" directory
# For example, running this (by clicking run or pressing Shift+Enter) will list all files under the input directory

import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

# You can write up to 20GB to the current directory (/kaggle/working/) that gets preserved as output when you create a version using "Save & Run All"
# You can also write temporary files to /kaggle/temp/, but they won't be saved outside of the current session
```

```
/kaggle/input/car-price-prediction/CarPrice_Assignment.csv
/kaggle/input/car-price-prediction/Data Dictionary - carprices.xlsx
```

## Load Data

In [2]:

```
df= pd.read_csv('/kaggle/input/car-price-prediction/CarPrice_Assignment.csv')
df.head()
```

Out[2]:

	car_ID	symboling	CarName	fueltype	aspiration	doornumber	carbody	drivewheel	engine	location	wheelbase	...	...
0	1	3	alfa-romero giulia	gas	std	two	convertible	rwd	front	88.6	...	...	...
1	2	3	alfa-romero stelvio	gas	std	two	convertible	rwd	front	88.6	...	...	...
2	3	1	alfa-romero Quadrifoglio	gas	std	two	hatchback	rwd	front	94.5	...	...	...
3	4	2	audi 100 ls	gas	std	four	sedan	fwd	front	99.8	...	...	...
4	5	2	audi 100ls	gas	std	four	sedan	4wd	front	99.4	...	...	...

5 rows x 26 columns



In [3]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
 #   Column              Non-Null Count  Dtype
---  -
 0   car_ID              205 non-null    int64
```

	car_id	year	non-null	dtype
1	symboling	205	non-null	int64
2	CarName	205	non-null	object
3	fueltype	205	non-null	object
4	aspiration	205	non-null	object
5	doornumber	205	non-null	object
6	carbody	205	non-null	object
7	drivewheel	205	non-null	object
8	enginelocation	205	non-null	object
9	wheelbase	205	non-null	float64
10	carlength	205	non-null	float64
11	carwidth	205	non-null	float64
12	carheight	205	non-null	float64
13	curbweight	205	non-null	int64
14	enginetype	205	non-null	object
15	cylindernumber	205	non-null	object
16	enginesize	205	non-null	int64
17	fuelsystem	205	non-null	object
18	boreratio	205	non-null	float64
19	stroke	205	non-null	float64
20	compressionratio	205	non-null	float64
21	horsepower	205	non-null	int64
22	peakrpm	205	non-null	int64
23	citympg	205	non-null	int64
24	highwaympg	205	non-null	int64
25	price	205	non-null	float64

In [4]:

Out[4]:

In [5]:

Out[5]:

mean	103.000000	0.834146	98.756585	174.049268	65.907805	53.724878	2555.565854	126.907317	3.329756	3.255
std	59.322565	1.245307	6.021776	12.337289	2.145204	2.443522	520.680204	41.642693	0.270844	0.313
min	1.000000	-2.000000	86.600000	141.100000	60.300000	47.800000	1488.000000	61.000000	2.540000	2.070
25%	52.000000	0.000000	94.500000	166.300000	64.100000	52.000000	2145.000000	97.000000	3.150000	3.110
50%	103.000000	1.000000	97.000000	173.200000	65.500000	54.100000	2414.000000	120.000000	3.310000	3.290
75%	154.000000	2.000000	102.400000	183.100000	66.900000	55.500000	2935.000000	141.000000	3.580000	3.410
max	205.000000	3.000000	120.900000	208.100000	72.300000	59.800000	4066.000000	326.000000	3.940000	4.170

In [6]:

```
df.nunique()
```

Out[6]:

```
car_ID          205
symboling        6
CarName         147
fueltype         2
aspiration       2
doornumber       2
carbody          5
drivewheel       3
enginelocation   2
wheelbase       53
carlength       75
carwidth        44
carheight       49
curbweight      171
enginetype       7
cylindernumber   7
enginesize      44
fuelsystem       8
boreratio       38
stroke          37
compressionratio 32
horsepower      59
peakrpm         23
citympg         29
highwaympg      30
price          189
dtype: int64
```

In [7]:

```
df['Company']= df['CarName'].apply(lambda x: x.split(" ")[0])
df['Company']= df['Company'].apply(lambda x: x.lower())
```

In [8]:

```
df.head()
```

Out[8]:

	car_ID	symboling	CarName	fueltype	aspiration	doornumber	carbody	drivewheel	enginelocation	wheelbase	...	f
0	1	3	alfa-romero giulia	gas	std	two	convertible	rwd	front	88.6	...	f
1	2	3	alfa-romero stelvio	gas	std	two	convertible	rwd	front	88.6	...	f
2	3	1	alfa-romero Quadrifoglio	gas	std	two	hatchback	rwd	front	94.5	...	f
3	4	2	audi 100 ls	gas	std	four	sedan	fwd	front	99.8	...	f
4	5	2	audi 100ls	gas	std	four	sedan	4wd	front	99.4	...	f

5 rows x 27 columns

In [9]:

```
df['Company'].unique()
```

Out[9]:

```
array(['alfa-romero', 'audi', 'bmw', 'chevrolet', 'dodge', 'honda',  
      'isuzu', 'jaguar', 'maxda', 'mazda', 'buick', 'mercury',  
      'mitsubishi', 'nissan', 'peugeot', 'plymouth', 'porsche',  
      'porcshce', 'renault', 'saab', 'subaru', 'toyota', 'toyouta',  
      'vokswagen', 'volkswagen', 'vw', 'volvo'], dtype=object)
```

In [10]:

```
df['Company'].replace({'maxda': 'mazda', 'porcshce': 'porsche', 'toyouta': 'toyota', 'voks  
wagen': 'volkswagen', 'vw': 'volkswagen'}, inplace=True)
```

In [11]:

```
df['Company'].unique()
```

Out[11]:

```
array(['alfa-romero', 'audi', 'bmw', 'chevrolet', 'dodge', 'honda',  
      'isuzu', 'jaguar', 'mazda', 'buick', 'mercury', 'mitsubishi',  
      'nissan', 'peugeot', 'plymouth', 'porsche', 'renault', 'saab',  
      'subaru', 'toyota', 'volkswagen', 'volvo'], dtype=object)
```

In [12]:

```
df.drop(['car_ID', 'CarName'], axis=1, inplace=True)
```

In [13]:

```
cat_cols= [col for col in df.columns if df[col].dtype=='object']  
num_cols= [col for col in df.columns if df[col].dtype!='object']
```

In [14]:

```
cat_cols
```

Out[14]:

```
['fueltype',  
 'aspiration',  
 'doornumber',  
 'carbody',  
 'drivewheel',  
 'enginelocation',  
 'enginetype',  
 'cylindernumber',  
 'fuelsystem',  
 'Company']
```

In [15]:

```
num_cols  
num_cols.remove('price')  
num_cols
```

Out[15]:

```
['symboling',  
 'wheelbase',  
 'carlength',  
 'carwidth',  
 'carheight',  
 'curbweight',  
 'enginesize',  
 'boretostratio',  
 'stroke',  
 'compressionratio',  
 'displacement',  
 'horsepower',  
 'peakrpm',  
 'citympg',  
 'highwaympg',  
 'mpg']
```







```
'stroke',
'compressionratio',
'horsepower',
'peakrpm',
'citympg',
'highwaympg']
```







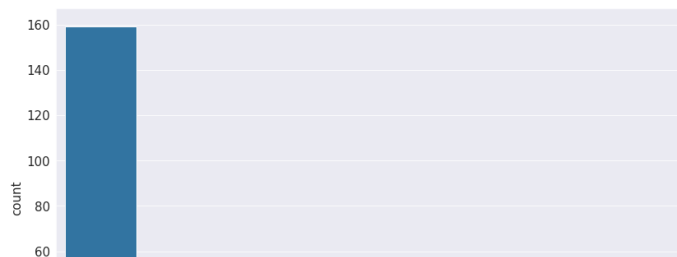





## Exploratory Data Analysis



In [16]:

```
i=1
plt.figure(figsize=(30,100))
for col in cat_cols:
    plt.subplot(10,2,i)
    sns.countplot(df[col])
    plt.xticks(rotation=90, fontsize=15)
    plt.yticks(fontsize=15)
    plt.xlabel(col, fontsize=15)
    plt.ylabel('count', fontsize=15)

    i+=1
    plt.subplot(10,2,i)
    sns.boxplot(x=df[col], y=df['price'])
    plt.xticks(rotation=90, fontsize=15)
    plt.yticks(fontsize=15)
    plt.xlabel(col, fontsize=15)
    plt.ylabel('price', fontsize=15)
    i+=1
plt.show()
```







## Observations




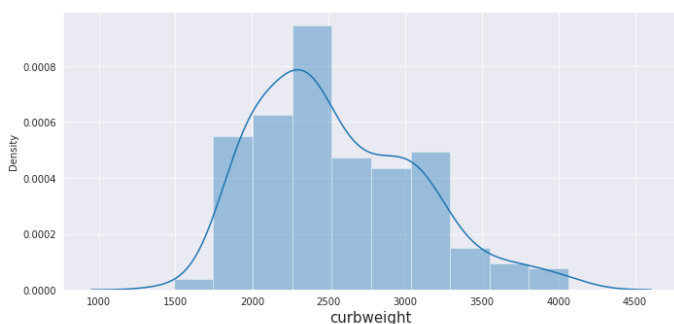

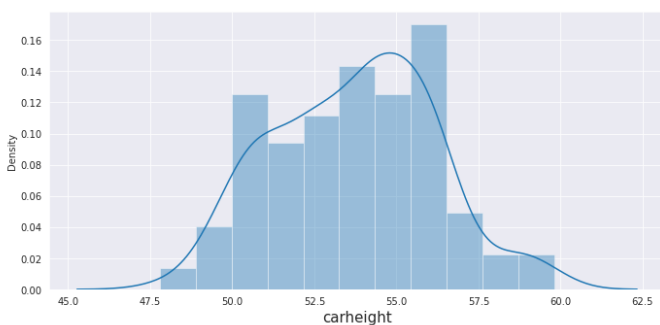

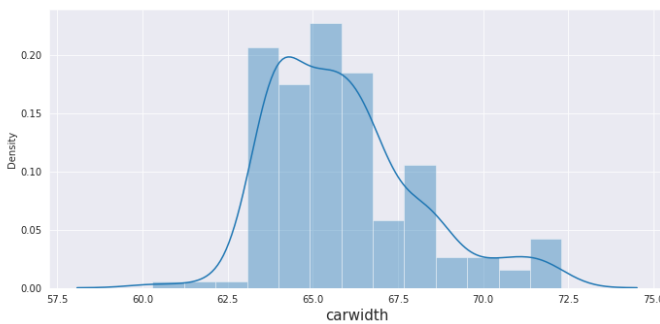






1. Diesel cars tend to be slightly higher-priced than gas cars
2. The number of gas cars is significantly higher than diesel cars
3. Turbo cars are higher priced than std cars
4. Sedan and hatchbacks account for more than 75% of total cars
5. Real wheel drive cars are higher priced than other drive cars
6. The median price of cars with engine at back is over 3 times the median price of cars with engine in front
7. Majority of the cars have ohc engines but, ohcv engines seem to be higher priced
8. General trend is that the price of car increase with increase in number of cylinders
9. Toyota is the most liked car
10. Porsche, Jaguar, BMW, Buick are high range cars

In [17]:















```
i=1
plt.figure(figsize=(25,100))
for col in num_cols:
    plt.subplot(16,2,i)
    sns.distplot(df[col])
    plt.xlabel(col,fontsize=15)
    plt.xticks(fontsize=10)
    i+=1



plt.subplot(16,2,i)
sns.scatterplot(x=df[col], y=df['price'])
plt.xlabel(col,fontsize=15)
plt.xticks(fontsize=10)
plt.ylabel('price', fontsize=15)
```

```
plt.yticks(fontsize=10)
i+=1
```









# Observations

- 1. Positive co-relation between price and car-length,car width, curb weight, engine size, horsepower
- 2. Negative co-relation between price and city mileage, highway mileage
- 3. There is no relation of symboling with price, hence we drop the column

In [18]:

```
df.drop('symboling', axis=1, inplace=True)
num_cols.remove('symboling')
```

In [19]:

```
df.head()
```

Out[19]:

	fueltype	aspiration	doornumber	carbody	drivewheel	enginelocation	wheelbase	carlength	carwidth	carheight	...	fu
0	gas	std	two	convertible	rwd	front	88.6	168.8	64.1	48.8	...	
1	gas	std	two	convertible	rwd	front	88.6	168.8	64.1	48.8	...	
2	gas	std	two	hatchback	rwd	front	94.5	171.2	65.5	52.4	...	
3	gas	std	four	sedan	fwd	front	99.8	176.6	66.2	54.3	...	
4	gas	std	four	sedan	4wd	front	99.4	176.6	66.4	54.3	...	

5 rows x 24 columns



In [20]:

```
num_cols
```

Out[20]:

```
['wheelbase',
 'carlength',
 'carwidth',
 'carheight',
 'curbweight',
 'enginesize',
 'boreratio',
 'stroke',
 'compressionratio',
 'horsepower',
 'peakrpm',
 'citympg',
 'highwaympg']
```

# Preprocessing

In [21]:

```
from sklearn.preprocessing import LabelEncoder
le= LabelEncoder()

df[cat_cols]= df[cat_cols].apply(lambda x: le.fit_transform(x))
```

In [22]:

```
df[cat_cols].head()
```

Out[22]:

	fueltype	aspiration	doornumber	carbody	drivewheel	enginelocation	enginetype	cylindernumber	fuelsystem	Company
0	1	0	1	0	2	0	0	2	5	0
1	1	0	1	0	2	0	0	2	5	0
2	1	0	1	2	2	0	5	3	5	0
3	1	0	0	3	1	0	3	2	5	1
4	1	0	0	3	0	0	3	1	5	1

In [23]:

```
df.head()
```

Out[23]:

	fueltype	aspiration	doornumber	carbody	drivewheel	enginelocation	wheelbase	carlength	carwidth	carheight	...	fuels
0	1	0	1	0	2	0	88.6	168.8	64.1	48.8	...	
1	1	0	1	0	2	0	88.6	168.8	64.1	48.8	...	
2	1	0	1	2	2	0	94.5	171.2	65.5	52.4	...	
3	1	0	0	3	1	0	99.8	176.6	66.2	54.3	...	
4	1	0	0	3	0	0	99.4	176.6	66.4	54.3	...	

5 rows x 24 columns



In [24]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 24 columns):
#   Column              Non-Null Count  Dtype
---  -
0   fueltype            205 non-null   int64
1   aspiration          205 non-null   int64
2   doornumber          205 non-null   int64
3   carbody             205 non-null   int64
4   drivewheel          205 non-null   int64
5   enginelocation      205 non-null   int64
6   wheelbase           205 non-null   float64
7   carlength           205 non-null   float64
8   carwidth            205 non-null   float64
9   carheight           205 non-null   float64
10  curbweight          205 non-null   int64
11  enginetype          205 non-null   int64
12  cylindernumber      205 non-null   int64
13  enginesize          205 non-null   int64
14  fuelsystem          205 non-null   int64
15  boreratio           205 non-null   float64
16  stroke              205 non-null   float64
17  compressionratio    205 non-null   float64
18  horsepower          205 non-null   int64
19  neakrrm             205 non-null   int64
```

```

20    fueltype      205 non-null    int64
21    highwaympg    205 non-null    int64
22    price         205 non-null    float64
23    Company       205 non-null    int64

```

dtypes: float64(8), int64(16)


memory usage: 38.6 KB

In [25]:

```

plt.figure(figsize=(15,15))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm')
plt.show()

```



In [26]:

```

from sklearn.preprocessing import StandardScaler

ss=StandardScaler()
df[num_cols]= ss.fit_transform(df[num_cols])

```

In [27]:

```
df.head()
```

Out[27]:

	fueltype	aspiration	doornumber	carbody	drivewheel	enginelocation	wheelbase	carlength	carwidth	carheight	...	fuels
0	1	0	1	0	2	0	-1.690772	-0.426521	0.844782	-2.020417	...	
1	1	0	1	0	2	0	-1.690772	-0.426521	0.844782	-2.020417	...	
2	1	0	1	2	2	0	-0.708596	-0.231513	0.190566	-0.543527	...	
3	1	0	0	3	1	0	0.173698	0.207256	0.136542	0.235942	...	
4	1	0	0	3	0	0	0.107110	0.207256	0.230001	0.235942	...	

5 rows x 24 columns



## Training our Model

In [28]:

```
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from xgboost import XGBRegressor
from sklearn.metrics import mean_squared_error, r2_score
```

In [29]:

```
X= df.drop('price', axis=1)
y=df['price']

X_train,X_test,y_train,y_test= train_test_split(X,y,test_size=0.2, random_state= 42)
```

In [30]:

```
lr= LinearRegression()
lr.fit(X_train,y_train)
y_pred= lr.predict(X_test)

rmse= (mean_squared_error(y_test,y_pred))**(1/2)
r2= r2_score(y_test,y_pred)

print(rmse)
print(r2)
```

```
3483.207163635308
0.8463122094668135
```

In [31]:

```
dt= DecisionTreeRegressor(criterion='mse', splitter='best', max_depth=100, min_samples_le
af= 5, random_state=42)
dt.fit(X_train,y_train)
y_pred= dt.predict(X_test)

rmse_train= (mean_squared_error(dt.predict(X_train),y_train))**(1/2)
rmse= (mean_squared_error(y_test,y_pred))**(1/2)
r2= r2_score(y_test,y_pred)
print(rmse_train)
print(rmse)
print(r2)
```

```
1546.3852781211913
```

2708.767381032721  
0.9070553964342399

In [32]:

```
rf= RandomForestRegressor(max_depth=10, criterion='mse', min_samples_leaf=2, random_state=42, verbose=1)
rf.fit(X_train,y_train)
y_pred= rf.predict(X_test)
rmse_train= (mean_squared_error(rf.predict(X_train),y_train))**(1/2)
rmse= (mean_squared_error(y_test,y_pred))**(1/2)
r2= r2_score(y_test,y_pred)
print(rmse_train)
print(rmse)
print(r2)
```

1182.8418416395994  
1953.5505961315666  
0.9516573910332784

```
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 100 out of 100 | elapsed: 0.2s finished
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 100 out of 100 | elapsed: 0.0s finished
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 100 out of 100 | elapsed: 0.0s finished
```

In [33]:

```
xgb= XGBRegressor()
xgb.fit(X_train, y_train)
y_pred= xgb.predict(X_test)
rmse_train= (mean_squared_error(xgb.predict(X_train),y_train))**(1/2)
rmse= (mean_squared_error(y_test,y_pred))**(1/2)
r2= r2_score(y_test,y_pred)
print(rmse_train)
print(rmse)
print(r2)
```

283.4366075958827  
2472.360172889496  
0.9225708957261225

## Conclusion

**All our models are overfitting the data. Tried many different combinations with GridSearchCV but there is still overfitting. This is because the amount of data is very less. However, with the given data, RandomForestRegressor is the best fit**

In [ ]: